

Alexander Loftus

AI researcher & communicator with 7+ years of experience in deep learning & machine learning. Kaggle \$100k competition winner, forthcoming CUP textbook author, and organizer of a 200-person mechanistic interpretability conference. Research in code interpretability, attribution, and evaluation for LLMs. Seeking role where deep technical depth, public speaking, and storytelling ability can combine.

Career highlights:

Textbook author: Authored a [524-page technical book](#) on statistical network ML (Cambridge Univ. Press, Nov 2025)

Organizer, Teacher, & Communicator: Organized the [New England Mechanistic Interpretability](#) (NEMI) conference; YouTube lecture series creator; delivered 10+ invited talks to 20-300 attendees; taught hundreds of students through meetups, summer camps, and tutorials.

\$100k ML Competition winner: Part of a 4-person team that won 1st place in the Vesuvius Kaggle competition (1,249 teams); [featured on the cover of *Scientific American*](#).

High-impact Interpretability research: Subliminal learning work featured in [YouTube video](#) with 1m+ subscribers; best poster award at NeurIPS 2023 LatinX workshop.

Strategic Advisory Roles: [CBAI mentor](#) for Harvard/MIT students; advisor for [cybersecurity/mechanistic interpretability](#) startup.

Cloud & AI Infrastructure: First author on [ICLR paper](#) on scaling up AI systems for interpretability; AWS experience [scaling up an AI pipeline](#) for computational neuroscience

EDUCATION

Northeastern University

Boston, MA

PhD Student, Computer Science

2024-Present

Advisor: [Dr. David Bau](#)

Focus on mechanistic interpretability in code LLMs. Data attribution, representation learning, causality.

Johns Hopkins University

Baltimore, MD

MSE Biomedical Engineering: Machine Learning & Data Science Focus

2020-2022

Advisor: [Dr. Joshua Vogelstein](#)

Thesis: [Hands-On Network Machine Learning](#)

dean's list, highest honors, GPA 3.97/4.0.

Western Washington University

Bellingham, WA

BS Behavioral Neuroscience — *Minors:* Chemistry, Philosophy

2014-2018

Founder & President, Computational Neuroscience Club

Vice President, Neuroscience Club

Built computational neuroscience club from scratch, taught weekly seminars.

EXPERIENCE

Data Scientist

San Diego, CA

Creyon Bio

2022-2024

Large Protein Models For Splice-Site Prediction: Explored splice site prediction in LLMs trained on protein sequences.

Pre-training, fine-tuning, and benchmarking+evals.

ML for Toxicity Prediction: Developed a novel contrastive learning pipeline to predict oligo toxicity from 3-D electrostatic maps; increased classification AUC from 0.73 to 0.88.

Neuron Toxicity Detection: Developed scalable neuron segmentation and toxicology classification pipeline.

All Projects: Helped shape Series B narrative through presentations to C-suite.

Machine Learning Research Engineer

Rockville, MD

Blue Halo

2021-2022

Conditional Image Generation with Generative Adversarial Networks: Replaced GAN pipeline with diffusion-model synthetic data generator. Immediate 10x training run reliability.

Detecting Objects with Enhanced Yolo and Knowledge Graphs: Led knowledge graph effort for object detection project.

Delivered live demos to program officers.

Geometric Multi-Resolution Analysis: Led infra for document clustering & analysis method.

Research Software Engineer	Baltimore, MD
NeuroData Lab, Johns Hopkins University — Dr. Joshua Vogelstein	2018-2021
<i>MRI-to-Graphs</i> : Optimized a diffusion MRI pipeline with Kubernetes, Docker, and AWS Batch. Halved runtime and cut cloud costs by 40%.	
<i>Graspologic</i> : Co-authored an open-source graph statistics library. Later adopted by Microsoft Research for large-scale network analysis.	
Assistant Director	Seattle, WA
iD Tech Camps — University of Washington	2014-2018 summers
<i>Leader and Manager</i> : Managed 10+ instructors/week and 300+ students.	
<i>Curriculum Designer</i> : Authored game development curriculum deployed to 50+ locations, impacting 10k+ students nationwide.	

SKILLS SUMMARY

Languages: Python, Bash, R, Rust, SQL

Tools & Frameworks: docker, kubernetes, pytorch, pytorch-lightning, VLLM, AWS, google cloud (GCP), numpy, scipy, pandas, polars, sklearn, seaborn, matplotlib, photoshop, SQL, weights & biases, mlflow, linux, cursor, ray, claude code, codex-cli

Areas of Expertise: LLMs for code, interpretability, transformers, GPUs and CUDA, linear algebra, probability & statistics, deep learning, information theory, diffusion models, convolutional autoencoders, public speaking, leadership & management, teaching, natural language processing, computer vision

Soft Skills: Public Speaking, Technical Writing, Leadership, Mentorship, Community-Building, Confidence & Charisma

TEXTBOOK

Hands-On Network Machine Learning with Python: *Eric Bridgeford, Alexander R. Loftus, Joshua Vogelstein.* Cambridge University Press, in copy-editing phase. To be printed November 2025.

Spectral representation theory on networks. 524 pages, 147 figures.

SELECTED PUBLICATIONS

* indicates equal contribution.

🏆 indicates best poster.

Token Entanglement in Subliminal Learning: *A. Zur, Z. Ying, A.R. Loftus, et al.* NeurIPS mechanistic interpretability workshop 2025.

Investigation on token entanglement in LLMs. Featured in [Welch Labs video](#) on YouTube.

NNsight and NDIF: Democratizing Access to Open-Weight Foundation Model Internals: *A.R. Loftus*, J.Fiotto-Kaufman*, et al.* ICLR 2025.

Open-source suite for probing & manipulating LLM weights without engineering overhead. Ray GCS Service backend with AWS object storage and VLLM for inference speed.

🏆A Saliency-based Clustering Framework for Identifying Aberrant Predictions : *A. Tersol Montserrat, A.R. Loftus, Y. Daihes.* Paper, **NeurIPS** LatinX AI Workshop, 2023.

Detects spurious feature reliance via saliency embeddings.

A low-resource reliable pipeline to democratize multi-modal connectome estimation and analysis: *J. Chung, R. Lawrence, A.R. Loftus, et al.* Paper, in review at Nature Methods, 2024

Transforms diffusion MRI scans into graphs; open-sourced ([code](#))

LEADERSHIP & COMMUNITY ENGAGEMENT

Conference Organizer	NEMI
Running 200+ person interpretability conference; Raised \$17,000 grant funding.	2025
Research Mentor	CBAI
Mentoring Harvard/MIT students in Summer 2025	2025
Strategic Advisor	Krnel.ai
Advisor to cybersecurity-focused startup specializing in interpretability tooling for AI systems.	2025
Meetup Speaker	SDML
Speaker & organizer for San Diego AI Meetups.	2023–2024
Hackathon Organizer	NeuroData Workshop
Helped organize hackathon & workshop to explore statistics for high-dimensional testing.	2019

TALKS & DEMOS

White-Box Techniques for Code LLMs: Influence Benchmarking, the Attendome, and Variable State Debugging: *Lawrence Livermore National Laboratory, 2025*

Invited talk on interpretability for code LLMs.

A Shared Infrastructure for Interpretability: *FAR AI Tech. Innovations for AI Policy Conf., 2025*

Invited demo for DC policymakers; showcased live editing of GPT2 internals

State of the Art in Knowledge Editing: *A.R. Loftus, 2023*

Survey talk on LLM knowledge-editing methods.

1st Place Solution - Vesuvius Ink Competition: *R. Chesler, A.R. Loftus, A. Tersol Montserrat, T. Kyi, 2023*

Walkthrough of winning \$100,000 ink-detection model.

ICML Conference Highlights: *A.R. Loftus, 2023*

Selected breakthroughs from ICML. Presented to biotech execs and SDML meetup group.

Working with LLMs: *AI San Diego Conference, 2023.*

Invited talk: Introduction to LLM engineering. 300+ attendees

Linear Algebra, from Dot Products to Neural Networks: *A.R. Loftus, 2023.*

Created a YouTube tutorial series on the fundamentals of linear algebra for machine learning.

FELLOWSHIPS & AWARDS

First Place Winner

Kaggle Vesuvius Competition, \$100,000. 2023

Khoury Distinguished Fellowship

Northeastern University PhD fellowship. 2024

GCP Research Grant

\$5,000 grant for computational research. 2025

Best Poster Award

NeurIPS 2023 LatinX AI Workshop. 2023

Harvard AI Safety Technical Fellowship

Harvard fellowship for technical work in AI safety. 2025

AWS Research Grant

\$10,000 grant for computational research on cloud services. 2019

TEACHING

Head Teaching Assistant

Foundations of Computational Biology and Bioinformatics, *EN.BME.410/634* Johns Hopkins University

Spring 2021

Teaching Assistant

NeuroData Design II, *EN.BME.438/638* Johns Hopkins University

Spring 2020

Teaching Assistant

NeuroData Design I, *EN.BME.437/637* Johns Hopkins University

Fall 2019

Teaching Assistant

Introduction to Behavioral Neuroscience, *PSY.220* Western Washington University

Winter 2017

Curriculum Designer

Built curriculum used across 50 locations in the United States by tens of thousands of students. iD Tech Camps

Spring 2017

Instructor

Taught programming and game design to high school students. iD Tech Camps

2014-2018 summers

FUN

Gaming: Starcraft 2 grandmaster in high school, local tournament winner

Music: Fingerstyle guitarist; performed at open mic nights.

Dancing: Partner dance instructor and competition winner (Fusion, West Coast Swing, Zouk)